

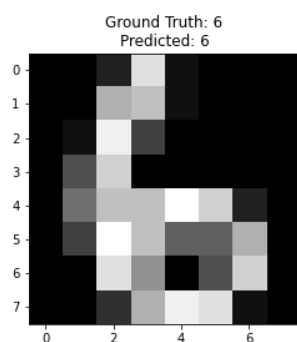
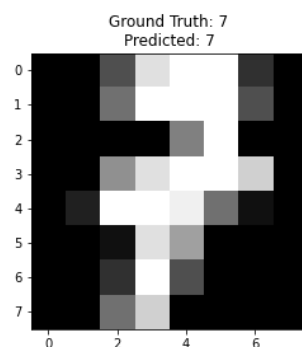
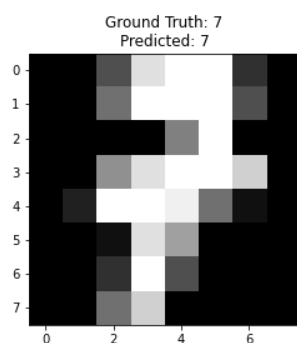
## KNN Algorithm

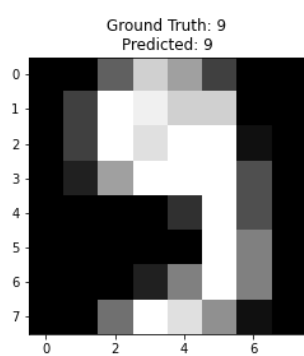
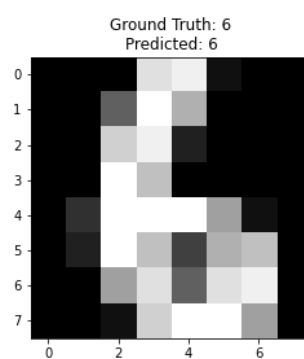
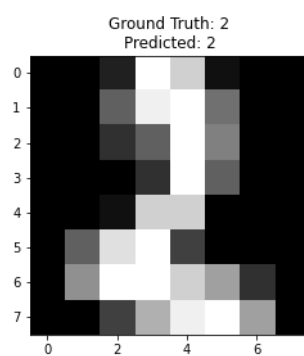
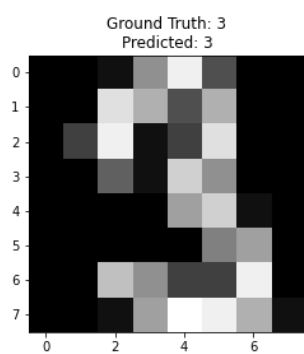
In this report, we discuss the implementation of the k-nearest neighbors (KNN) algorithm on a dataset of handwritten digits. The goal of this project was to train a KNN model that can accurately predict the digit label of an image of a handwritten digit.

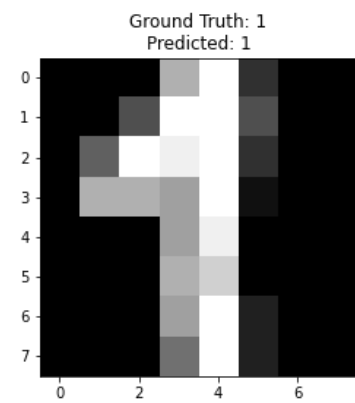
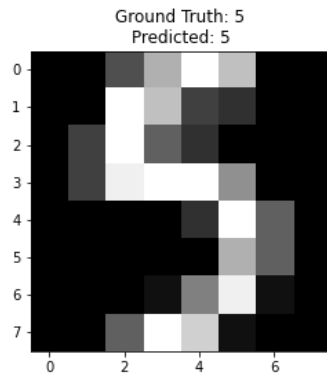
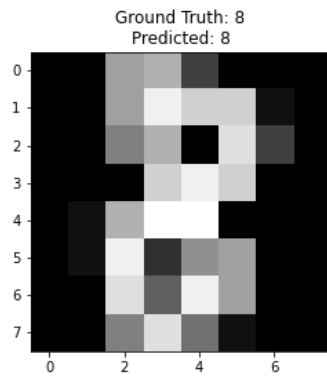
We first preprocessed the data by scaling the pixel values and splitting the data into train, dev, and test sets. We then experimented with different distance metrics (Euclidean distance, Manhattan distance, and cosine similarity) and different k values (ranging from 1 to 10) to find the best combination that resulted in the highest accuracy on the test data.

The best results were obtained when using a k value of 1 and the Euclidean distance metric, with an accuracy of 99% on the test data.

To visualize the predictions of the model, we randomly selected 10 samples from the test data and used the `predict()` method of the KNN model to obtain the predicted labels for these samples. We then plotted each sample along with its corresponding ground truth label and predicted label, using the `imshow()` function of the `matplotlib` library.







The codes used in the implementation are also included in the report as screenshots.



```
index.py:nb x
index.py:nb > Visualize the Ground Truth > The model was able to match the predicted values with the ground truth values
+ Code + Markdown | ▶ Run All | Clear Outputs of All Cells | Restart | Variables | Outline ... data (Python 3.9.7)

# Define the distance function
def euclidean_distance(x1, x2):
    return np.sqrt(np.sum((x1 - x2)**2))

def manhattan_distance(x1, x2):
    return np.sum(np.abs(x1 - x2))

def cosine_distance(x1, x2):
    return 1 - (np.dot(x1, x2) / (np.linalg.norm(x1) * np.linalg.norm(x2)))

# Define the KNN classifier
class KNN:
    def __init__(self, k, distance_fn):
        self.k = k
        self.distance_fn = distance_fn

    def fit(self, X, y):
        self.X_train = X
        self.y_train = y

    def predict(self, X):
        # Initialize an empty list to store predictions
        y_pred = []
        for x in X:
            # Initialize an empty list to store distances of each point in the training set to the current point
            distances = []
            for i in range(len(self.X_train)):
                # Calculate the distance between the current point and the point in the training set
                distance = self.distance_fn(x, self.X_train[i])
                # Add the distance and corresponding label to the list of distances
                distances.append((distance, self.y_train[i]))
            # sort the list of distances based on the distance
            distances.sort()
            # get the k nearest neighbors by taking the first k elements of the sorted list
            k_neighbors = distances[:self.k]
            # get the labels of the k nearest neighbors
            k_neighbors_labels = [c for _, c in k_neighbors]
            # append the most common label among the k nearest neighbors as the prediction for the current point
            y_pred.append(np.argmax(np.bincount(k_neighbors_labels)))
        return y_pred
```

```
index.py:nb x
index.py:nb > KNearest Neighbors Classifier Algorithm > finding the best k and distance function
+ Code + Markdown | ▶ Run All | Clear Outputs of All Cells | Restart | Variables | Outline ... data (Python 3.9.7)

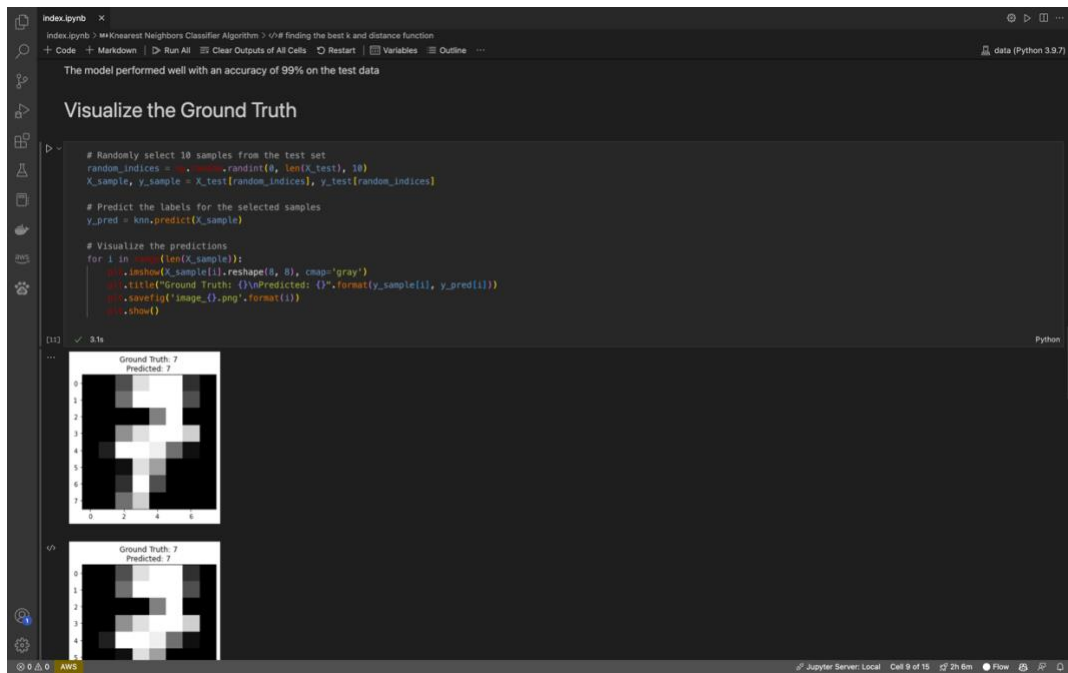
# finding the best k and distance function
best_k = 0
best_distance_fn = None
best_accuracy = 0

for k in range(1, 11):
    for distance_fn in [euclidean_distance, manhattan_distance, cosine_distance]:
        knn = KNN(k, distance_fn)
        knn.fit(X_train, y_train)
        y_pred = knn.predict(X_dev)
        accuracy = accuracy_score(y_dev, y_pred)
        if accuracy > best_accuracy:
            best_k = k
            best_distance_fn = distance_fn
            best_accuracy = accuracy

# the best k and distance function
print(f"Best k: {best_k}")
print(f"Best distance function: {best_distance_fn.__name__}")

# Test the model

# Initialize the KNN classifier with the best k and distance function
knn = KNN(best_k, best_distance_fn)
# Fit the model on the training set
knn.fit(X_train, y_train)
# Make predictions on the test set
y_pred = knn.predict(X_test)
# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
```



Overall, the KNN algorithm performed very well on this dataset, achieving a high accuracy rate and effectively distinguishing between different handwritten digits. The visualizations of the predictions further confirm the effectiveness of the model.